

# CS450

## Programming Assignment 4: Report

November 29, 2017

Ling Xing Zheng  
A20343483

Venkata Krishna Chaitanya  
Polavarapu  
A20378279

### Question 1

#### Extent based file system

The following files are modified to implement extent based file system.

<code>fcntl.h</code>	: Modified to include <code>#define O_EXTENT 0x020</code> , used in create function call.
<code>fs.c</code>	: Modified to implement changes for extent based file system.
<code>stat.h</code>	: Modified to include new file type ( <code>#define T_EXT_FILE 4</code> ). Changes made to " <code>struct stat</code> " to extend the functionality of <code>fstat()</code>
<code>sysfile.c</code>	: Changes made to <code>sys_open</code> to appropriately pass the file type during file creation.

#### Data allocation strategy

1. The idea is that we use existing inode pointers to store the (block address, length). The first 24 bits holds the address and the last 8 bits hold the length.
2. Whenever a data block of a file is needed, function `bmap` is called. This function maps the file block number to disk block number. Most of the changes are made here.
3. When ever `bmap()` is called, we first check if the requested file block is on disk. We do this by checking if the given block number is within the `[ip->addrs, ip->addr+len)`. Repeat the previous check for all the available extents in the inode. If found, we return the corresponding block address. The corresponding code is shown below (`bmap` from `fs.c`).

```

    if (ip->type == T_EXT_FILE){
387      // Check if the block is already in use
388      int i = 0;
389      uint currLen = 0;
391      while(ip->addrs[i]){
392          // Extract the last 8 bits, which gives the length
393          uint length = EXT_LEN(ip->addrs[i]);
394
395          // Check of the requested block is within the extent
396          if (bn >= currLen && bn < length + currLen){
398              addr = EXT_ADDR(ip->addrs[i]) + bn - currLen;
399              return addr;
400          }
401
402          // if not move to next extent
403          currLen += length;

```

4. If the block is not found, we ask for a new block. If the new block is contiguous with respect to the last extent in the inode, we increase its length. If not we add a new entry in the inode address table and return the disk address. If however the entries are  $\geq$  NDIRECT we throw a kernel panic(). The corresponding code of bmap from fs.c.

```

410     if(i >= NDIRECT){
411         panic("Extent file allocation exceeded\n");
412     }
414     uint newDiskBlock = balloc(ip->dev);
415     // Check if newAddr is contiguous w.r.t to the previous extent
416     uint len = EXT_LEN(ip->addrs[i-1]);
417     if( (newDiskBlock == EXT_ADDR(ip->addrs[i-1]) + len)
418         && EXT_LEN(ip->addrs[i-1]) < 0xff){
419         ip->addrs[i-1] += 1;
421     }
422     else{
423         // if not, add a new entry in the table
424         ip->addrs[i] = GEN_ADDR(newDiskBlock, 1);
426     }
427     return newDiskBlock;

```

## Data de-allocation

1. When a file is to be deleted, the system call `unlink` eventually calls `itrunc` which frees up the inodes (by calling `bfree`).
2. Our changes for data de-allocation happens in `itrunc()`. If the file type is `T_EXT_FILE`, we read each extent, for each block in the extent we call `bfree()`. The code below(`fromitrunc()` in `fs.c`) shows what we said above

```
// Code to clean Extent files
467  if (ip->type == T_EXT_FILE){
468      for(i = 0; i < NDIRECT; i++){
469          if(ip->addrs[i]){
470              uint len = EXT_LEN(ip->addrs[i]);
471              for(j = 0; j < len; j++){
472                  bfree(ip->dev, EXT_ADDR(ip->addrs[i]) + j);
474              }
475              ip->addrs[i] = 0;
476          }
477      }
479      return;
480  }
```

## Question 2

### `fstat()` Implementation details

The following files are changed for implementing changes to `fstat()`

```
fs.c : Modified to implement changes to stat()
stat.h: Changes made to "struct stat" to extend the functionality of fstat()
```

1. The `struct stat` is changed in the following way to hold additional details for extent based files.

```
6  struct stat {
7      short type;    // Type of file
8      int dev;       // File system's disk device
9      uint ino;      // Inode number
10     short nlink;   // Number of links to file
11     uint size;     // Size of file in bytes
12     struct extents_details{
13         uint addr:24, // Extent Pointers
14         len:8 ;      // Extent Length
15     }extents[12];
16 };
```

2. The `stati()` was changed to collect more information if the file is extent based.

```
505 // Copy stat information from inode.
506 // Caller must hold ip->lock.
507 void
508 stati(struct inode *ip, struct stat *st)
509 {
510     st->dev = ip->dev;
511     st->ino = ip->inum;
512     st->type = ip->type;
513     st->nlink = ip->nlink;
514     st->size = ip->size;
515
516     // Get more info for extent files
517     int i = 0;
518     if (st->type == T_EXT_FILE){
519         for(i = 0; i < NDIRECT; i++)
520         {
521             st->extents[i].addr = EXT_ADDR(ip->addrs[i]);
522             st->extents[i].len = EXT_LEN(ip->addrs[i]);
523         }
524     }
525 }
```

## Testing Extent files/fstat()

We write the user code `assign4.c` that calls various function including `create`, `write`. We also use existing utilities like `wc`, `rm` for our testing purposes.

### Test Case 1:

The default files in xv6 system can have a maximum size of 70 KB. However extent based files hold file sizes  $> 70$  KB. We test this by writing 80KB to the file.

### Test Case 2:

We ensure that the default file system is not affected by our changes by writing to 2 files each of 1KB.

### Test Case 3:

We test our modified `fstat()` on extent based files to if it works as expected.

## Results of Testcase 1,2,3:

```
$ assign4
lseek failed, returned value: -1
lseek successful, offset is: 16
lseek successful, offset is: 700
lseek failed, returned value: -1
Wrote 80 * 1024 bytes successfully
File size 81920 bytes
Extent addr:606. Extent Len:102
Extent addr:710. Extent Len:20
Extent addr:732. Extent Len:8
Extent addr:742. Extent Len:2
Extent addr:746. Extent Len:10
Extent addr:759. Extent Len:10
Extent addr:771. Extent Len:8
```

Figure 1: output of assign4.c: We see the output written successfully. `fstat` displaying values correctly.

```

$ ls
. 1 1 512
.. 1 1 512
README 2 2 2290
assign4 2 3 16464
cat 2 4 13604
echo 2 5 12680
forktest 2 6 8400
grep 2 7 15424
init 2 8 13264
kill 2 9 12728
ln 2 10 12628
ls 2 11 14960
mkdir 2 12 12752
rm 2 13 12732
sh 2 14 23364
stressfs 2 15 13400
usertests 2 16 56276
wc 2 17 14248
zombie 2 18 12460
console 3 19 0
test.txt 4 20 81920
test_normal_1 2 21 1024
test_normal_2 2 22 1024
test_lseek_1 2 23 1024
test_lseek_2 2 24 1024
test_lseek_3 2 25 1212
test_lseek_4 2 26 1024

```

Figure 2: `ls` output showing that our file is indeed 80 KB

```
$ wc test.txt
0 1 81920 test.txt
$ wc test_normal_1
0 1 1024 test_normal_1
$ wc test_normal_2
0 1 1024 test_normal_2
```

Figure 3: `wc` output showing that the existing file/ extent file read writes work as expected